

Kinetic Heap Analysis

1. Abstract

A kinetic data structure is a data structure used to track an attribute of a system that is changing continuously as a function of time. So the system has some values, and for each value v , it is known that $v=f(t)$. Kinetic data structures allow queries on a system at the current virtual time t and two additional operations:

advance(t) - Advances the system to time.

change($v, f(t)$)- Alters the trajectory of value.

Additional operations may be supported. For example, kinetic data structures are often used with a set of points. In this case, the structure typically allows points to be inserted and deleted. In principle, this can be approximated by sampling the position of the points at fixed intervals of time and deleting and re-inserting each point into a “static” (traditional) data structure. However, such an approach is vulnerable to oversampling or under-sampling, depending on what interval of time is used, and can also be wasteful of computational resources.

Kinetic data structures are frequently used to calculate dynamic geo/spatial configurations like moving points over specific trajectories. However, it applies to everything that can be treated as mutating over time, a classical example is a priority queue with priorities as a function of time.

2. Introduction

There are many scientific researches dedicated to kinetic data structures. However, there are a limited number of practical implementations that could serve understanding and learning purposes. This article intends to partially fill the gap and to provide a proof of concept of “kinetic heap” implementation.

This approach includes the following steps:

- Review of the certificate-based approach.
- Creating a model for Kinetic Heap.
- Java code design and implementation.
- Basic performance tests.

What is not included:

- Advanced JMH performance tests, because of the requirement for a very thorough input data preparation. This is outside the scope of this article.
- The provided code could be further optimized as we may have alternatives in using relevant data structures (certificate Heap vs Binary Search Tree, lazy certificates invalidation vs immediate removal from the heap, etc. . .).
- Theoretical $O(N)$ notation complexity analysis.
- Observations of other kinetic structures like kinetic tournament, kinetic convex hull, kinetic sorted list, etc. . .

3. Related Works

It is worthy to mention several works:

- Leonidas J Guibas “Kinetic Data Structures”: <https://graphics.stanford.edu/courses/cs268-16-fall/Notes/g-kds.pdf>
It provides a good theoretical problem background for the subject.
- Fairly good Wiki article: https://en.wikipedia.org/wiki/Kinetic_data_structure
- Zahed Rahmati “Simple, Faster Kinetic Data Structures”: <https://dspace.library.uvic.ca/server/api/core/bitstreams/fc836d25-2016-4d68-b92e-003b95ef608d/content>

It is recommended to read it for better understanding of KDS in geo/space applications.

4. Problem Background

Heap is the basic collection structure to effectively perform the following operations:

- **getMin/Max()** - Return a minimum/maximum value.

- **insert(x)** - Insert a value.

The time complexity is $O(N \log(N))$ because the heap is a complete binary tree with the maximum height $\log(k)$ where k is the number of elements. It allows duplicates and usually is used as priority queues. The relevant Java class is **PriorityQueue**.

However, what will happen if priorities are mutating as a function of time? We will need to add the function:

- **advance(t)** - Advances the system to time t .

There are a lot of examples, starting from a set of objects having various initial temperatures at 0 point of time and heating with different rates. Generalizing, the function of a temperature from time could be a non-linear function. Alternatively, those could be space coordinates of neighboring points that can intersect at some moment, so priorities can swap for them at some moment.

A naive approach would be to remove all elements from the heap and insert all of them with new priorities at the current moment, thus we will preserve the heap invariant.

An obvious disadvantage of that solution is an increasing time complexity of **advance(t)** operation - $O(N^2)$

We need to design a structure that could mitigate this complexity and be quite close to the original logarithmic complexity.

5. Architecture of the Application and Implementation of the Solution

Prerequisites:

- We will be creating min heap.
- For simplicity, we will be using a discrete timescale.
- We will use a simple kinetic structure - priorities will be the linear function of time and initial temperature: $t(x) = t(0) + t * a$, where a is the temperature rate, $t(0)$ - an initial temperature. Thus, we will guarantee that a given pair of temperature lines will intersect only once (if converging to an intersection point). This point is easy to calculate and exclude from calculations once the intersection point is processed.
- We will be using a well-known certificate-based approach. Wiki describes this quite well: https://en.wikipedia.org/wiki/Kinetic_heap see (Implementation and operations, Dealing with certificate failures section). Certificates are just additional components allowing detection of elements that must be swapped with its root element at a given moment of time and helping to avoid full heap scan.




Figure 1: img.png

One can see which certificates is necessary to update when the condition $A > B$ fails. This is called a certificate failure (max heap example).

Implementation steps:

- Let's start from interfaces

Standard heap:

```
public interface IHeap<T extends Comparable<T>> {  
  
    T extractMin();  
  
    T getMin();  
  
    void insert(T data);  
  
    int size();  
  
    void clear();  
}
```

Kinetic heap:

```
public interface IKineticHeap extends IHeap<KineticElement> {  
  
    void fastForward(int nextTime);  
  
    int getCurTime();  
}
```

Kinetic element. The important attributes are:

id - element ID

initialPriority - initial priority

rate - temperature rate over time

```
public class KineticElement implements Comparable<KineticElement> {  
  
    private final int id;  
  
    private final double initialPriority;  
  
    private final double rate;  
    ...  
  
    public double getPriority() {  
        return initialPriority + rate * timeSupplier.get();  
    }  
    ...  
    public double getIntersectionTime(KineticElement other) {  
        if (rate - other.rate == 0) {  
            return Double.NEGATIVE_INFINITY;  
        }  
        return (other.initialPriority - initialPriority) / (rate - other.rate);  
    }  
  
    @Override  
    public int compareTo(KineticElement other) {  
        return Double.compare(getPriority(), other.getPriority());  
    }  
    ...  
}
```

Note that it is quite easy to calculate intersection time.

Certificate:

elementIdx - the reference index of a kinetic element

expirationTime - a certificate expiration time

ownIdx - an index of an element in a heap

```
public class Certificate implements Comparable<Certificate> {
    ...
    private final int elementIdx;

    private final double expirationTime;

    @Setter
    private int ownIdx = -1;

    @Override
    public int compareTo(Certificate other) {
        return Double.compare(expirationTime, other.expirationTime);
    }
    ...
}
```

Note that the heap is an array-based structure, so we will be using index-based access for certificates and elements. Therefore, to effectively remove certificates from the middle of the heap, we need to keep relevant indexes.

To make this synchronization easier, the helper `IEventSink` interface is created

```
public interface IEventSink<T extends Comparable<T>> {

    void onBubbleUpEventBeforeSwap(IHeap<T> heap, int idx, int parentIdx);

    void onBubbleUpEventAfterSwap(IHeap<T> heap, int idx, int parentIdx);

    void onBubbleUpEventNoChange(IHeap<T> heap, int idx);

    void onBubbleDownEventBeforeSwap(IHeap<T> heap, int idx, int parentIdx);

    void onBubbleDownEventAfterSwap(IHeap<T> heap, int idx, int parentIdx);

    void onBubbleDownEventNoChange(IHeap<T> heap, int idx);
}
```

It will help to intercept events on bubble up and bubble down events on inserting/removing elements from the heap.

- Concrete implementation:

Further, we will create a standard heap implementation what is of no real interest, see widespread web resources explaining how it works, for example:

<https://trykv.medium.com/algorithms-on-graphs-the-importance-of-heaps-e3e1385ae534>

Standard heap implementation `Heap` class:

<https://github.com/victor6567758/kinetic/blob/a90490c1196a8503c81faeef859628aa0ec9ce20/src/main/java/org/kinetic/heap/Heap.java>

The core logic of kinetic approach is implemented in `KineticHeap`:

<https://github.com/victor6567758/kinetic/blob/a90490c1196a8503c81faeef859628aa0ec9ce20/src/main/java/org/kinetic/heap/KineticHeap.java>

Those are the key elements:

heap - just the standard heap min heap, contains `KineticElement` objects.

certificates - the min heap containing certificates, so one can extract the soonest certificates which will fail.

curTime - this is the current time, it is modified per each call of `fastForward` (see `advance(x)`).

Note, that each heap element contains a certificate describing the relation between itself and the root element.

```
Certificate cert = heap.getValue(idx).getCertificate();
```

```

public class KineticHeap implements IKineticHeap {

    private final Heap<KineticElement> heap = new Heap<>(null);

    private final Heap<Certificate> certificates = new Heap<>(new CertificateEventSink());

    private int curTime;
    ...
}

```

Let's review `fastForward` function: It extracts the nearest certificate, then it invalidates up to 5 certificates around, make relevant elements swap (we have monotonically increasing priorities, so the child bubbles down), then inserts up to 4 certificates, as we don't need to update the certificate pointing the parent element (we are already done with this). Note that we use a current certificate time when inserting new ones to avoid generation of already processed certificates. This logic, probably, can be done better as to avoid possible edge conditions.

```

@Override
public void fastForward(int nextTime) {
    if (nextTime <= curTime) {
        return;
    }

    curTime = nextTime;

    while (true) {

        Certificate minCertificate = certificates.getMin();
        if (minCertificate == null || certificates.getMin().getExpirationTime() > nextTime) {
            break;
        }

        Certificate certificate = certificates.getMin();

        if (certificate.getOwnIdx() == -1) {
            throw new IllegalArgumentException();
        }

        int elemIdx = certificate.getElementIdx();
        int parentIdx = Heap.getParent(elemIdx);
        invalidateCertificates(elemIdx, parentIdx);

        heap.swap(elemIdx, parentIdx);

        insertCertificates(elemIdx, certificate.getExpirationTime());

    }
}

```

Let's review bubble up/down functions. This is quite a standard implementation, we just care about certificates invalidation if this is a non-root element (the root can't keep a certificate) and inserting new certificates if the one we are processing exceeds `curTime`. If no swap occurred, we are just inserting the only one certificate for the current index, because we don't need to invalidate other ones.

```

private int heapUp() {
    int curIndex = heap.size() - 1;
    while (curIndex > Heap.getRoot()) {
        int parentIndex = Heap.getParent(curIndex);
        if (heap.getValue(curIndex).compareTo(heap.getValue(parentIndex)) < 0) {
            if (curIndex != 0) {
                invalidateCertificates(curIndex, parentIndex);
            }
        }
    }
}

```

```

        heap.swap(curIndex, parentIndex);
        insertCertificates(curIndex, curTime);
    } else {
        break;
    }

    curIndex = parentIndex;
}

if (curIndex == heap.size() - 1) {
    createAndMaybeAddCertificate(curIndex, curTime);
}
return curIndex;
}

```

That's also a standard implementation. If no swap occurred, we don't do anything as we always start from the root.

```

private int heapDown() {
    int curIndex = 0;
    int size = heap.size();
    while (true) {
        int leftChildIndex = Heap.getLeftChild(curIndex);
        int rightChildIndex = Heap.getRightChild(curIndex);
        if (leftChildIndex >= size) {
            break;
        }

        boolean hasRight = rightChildIndex < size;
        int smallestChildIndex =
            hasRight && heap.getValue(rightChildIndex).compareTo(heap.getValue(leftChildIndex)) < 0
                ? rightChildIndex : leftChildIndex;

        if (heap.getValue(smallestChildIndex).compareTo(heap.getValue(curIndex)) < 0) {
            invalidateCertificates(smallestChildIndex, curIndex);
            heap.swap(smallestChildIndex, curIndex);
            insertCertificates(smallestChildIndex, curTime);
        } else {
            break;
        }
        curIndex = smallestChildIndex;
    }

    return curIndex;
}

```

Let's look at `extractMin/insert` methods. Because we do not remove from the middle, there is bubbling up or down. Remember to invalidate a relevant certificate before removing the last element.

```

@Override
public KineticElement extractMin() {
    KineticElement minElement = getMin();

    if (minElement != null) {
        int lastIdx = heap.size() - 1;

        heap.getValue(lastIdx).invalidateCertificate(certificates);
        KineticElement old = heap.setValue(heap.getValue(lastIdx), 0);
        heap.remove(heap.size() - 1);

        heapDown();
        return old;
    }
}

```

```

    return null;
}

@Override
public void insert(KineticElement data) {
    if (data == null) {
        throw new IllegalArgumentException("Invalid data");
    }
    heap.appendValue(data);
    heapUp();
}

```

The code what keeps the certificates heap invariants and makes it self-referencing correct indexes:

```

private final Heap<Certificate> certificates = new Heap<>(new CertificateEventSink());
...

private class CertificateEventSink implements IEventSink<Certificate> {

    @Override
    public void onBubbleUpEventBeforeSwap(IHeap<Certificate> heap, int idx, int parentIdx) {

    }

    @Override
    public void onBubbleUpEventAfterSwap(IHeap<Certificate> heap, int idx, int parentIdx) {
        setCertificateIndex(idx);
        setCertificateIndex(parentIdx);
    }

    @Override
    public void onBubbleDownEventBeforeSwap(IHeap<Certificate> heap, int idx, int parentIdx) {

    }

    @Override
    public void onBubbleDownEventAfterSwap(IHeap<Certificate> heap, int idx, int parentIdx) {
        setCertificateIndex(idx);
        setCertificateIndex(parentIdx);
    }

    @Override
    public void onBubbleUpEventNoChange(IHeap<Certificate> heap, int idx) {
        setCertificateIndex(idx);
    }

    @Override
    public void onBubbleDownEventNoChange(IHeap<Certificate> heap, int idx) {
        setCertificateIndex(idx);
    }
}
...

private void setCertificateIndex(int idx) {
    if (idx < certificates.size()) {
        Certificate certificate = certificates.getValue(idx);
        certificate.setOwnIdx(idx);
    }
}

```

```

    ...
}

```

6. Performance Evaluation

The detailed performance testing is outside the scope of this article.

It is essential to keep in mind those metrics must be observed:

- Responsiveness: A kinetic heap is responsive, since each certificate failure causes impacted keys to be swapped and leads to only a few certificates being replaced in the worst case.
- Locality: Each node is present in one certificate each along with its parent node and two child nodes (if present), meaning that each node can be involved in a total of three scheduled events in the worst case, thus kinetic heaps are local.
- Compactness: Each edge in the heap corresponds to exactly one scheduled event, therefore the number of scheduled events is exactly $n-1$, where n is the number of nodes in the kinetic heap. Thus, kinetic heaps are compact.
- Efficiency: Note that as per https://en.wikipedia.org/wiki/Kinetic_heap (Analysis of efficiency) the analysis is quite complex, but it is quite efficient when the number of certificates failures is rather small comparing to the number of elements. This condition is not fulfilled in my tests as I generated elements randomly. However, we still see performance gains compared to the trivial implementation: The micro-benchmarking was performed with JMH library, `SingleShotTime` mode because of the tests were quite computationally heavy.

Benchmark	(maxTimeSteps)	(n)	Mode	Score	Units
kineticHeapAddTimeForward	100	10	ss	0.146	ms/op
kineticHeapAddTimeForward	100	100	ss	0.383	ms/op
kineticHeapAddTimeForward	100	1000	ss	1.764	ms/op
kineticHeapAddTimeForward	100	10000	ss	8.198	ms/op
kineticHeapInserts	100	10	ss	0.075	ms/op
kineticHeapInserts	100	100	ss	0.229	ms/op
kineticHeapInserts	100	1000	ss	0.837	ms/op
kineticHeapInserts	100	10000	ss	4.618	ms/op
kineticHeapRemoves	100	10	ss	0.215	ms/op
kineticHeapRemoves	100	100	ss	0.483	ms/op
kineticHeapRemoves	100	1000	ss	7.558	ms/op
kineticHeapRemoves	100	10000	ss	50.519	ms/op
trivialHeapAddTimeForward	100	10	ss	1.233	ms/op
trivialHeapAddTimeForward	100	100	ss	1.666	ms/op
trivialHeapAddTimeForward	100	1000	ss	3.034	ms/op
trivialHeapAddTimeForward	100	10000	ss	18.022	ms/op
trivialHeapInserts	100	10	ss	0.025	ms/op
trivialHeapInserts	100	100	ss	0.046	ms/op
trivialHeapInserts	100	1000	ss	0.186	ms/op
trivialHeapInserts	100	10000	ss	0.850	ms/op
trivialHeapRemoves	100	10	ss	0.035	ms/op
trivialHeapRemoves	100	100	ss	0.245	ms/op
trivialHeapRemoves	100	1000	ss	0.363	ms/op
trivialHeapRemoves	100	10000	ss	2.379	ms/op

You can look at `testExample` folder to see a visual representation of the heap and what happens with certificates when we are forwarding time.

Compare `trivialHeapAddTimeForward` vs `kineticHeapAddTimeForward`

6. References

- <https://medium.com/@hasithalgamge/unit-testing-part-4-performance-based-unit-test-af83ce6a3966>
- <https://trykv.medium.com/algorithms-on-graphs-the-importance-of-heaps-e3e1385ae534>
- <https://medium.com/@gavinleegoodship23/implementing-kinetic-data-structures-in-python-managing-real-time-data-c6b107c4debc>
- https://en.wikipedia.org/wiki/Kinetic_heap

- <https://graphics.stanford.edu/courses/cs268-16-fall/Notes/g-kds.pdf>
- <https://github.com/6851-2021/kinetic>
- <https://github.com/frankfarrell/kds4j>
- <https://dspace.library.uvic.ca/server/api/core/bitstreams/fc836d25-2016-4d68-b92e-003b95ef608d/content>
- <https://www.baeldung.com/cs/heap-vs-binary-search-tree#:~:text=The%20Heap%20differs%20from%20a,as%20an%20array%20>